

Senior Project - Exploratory Data Analysis Using Python

April 29, 2018

1 e-Commerce Exploratory Data Analysis Using Python

1.0.1 Faculty Advisor: Dr. Patrick King

Date: Tuesday, April 3rd, 2018 This is a project of doing an exploratory data analysis (EDA) in Python Jupyter Notebook. This project is part II of the first phase, which is a comparison of the EDA process using R Markdown. We will run a parallel analysis and compare the two softwares while providing essential background information on the advantages and capabilities of the two softwares.

First, let's import the packages needed to create the plots we will use to describe our data as well as making hypotheses. We start with the **pandas**, **numpy**, and **matplotlib** packages which are very useful for data exploration.

```
In [40]: # import python packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pandas import ExcelWriter
from pandas import ExcelFile
```

We will use the `read_excel` module from **pandas** package to import the Excel spreadsheet containing the Zappos customer transactions. We can take a look at the columns and the type of data each holds. We can also identify the number of missing values for each columns as well as their range if they are type *int*.

```
In [41]: mydata = pd.read_excel('C:/Users/Dorcas/OneDrive - University of Houston Downtown/Sen
print("Column headings: ")
# List all column names
print(mydata.columns)
```

Column headings:

```
Index(['day', 'site', 'new_customer', 'platform', 'visits',
       'distinct_sessions', 'orders', 'gross_sales', 'bounces', 'add_to_cart',
       'product_page_views', 'search_page_views', 'conversion_rate',
       'bounce_rate', 'add_to_cart_rate'],
      dtype='object')
```

In [42]: mydata.head(25)

```
Out[42]:
```

	day	site	new_customer	platform	visits	distinct_sessions	\
0	2013-01-01	Acme	1.0	Android	24	16	
1	2013-01-01	Acme	1.0	BlackBerry	0	0	
2	2013-01-01	Sortly	1.0	iPad	0	0	
3	2013-01-01	Acme	1.0	Windows	922	520	
4	2013-01-01	Botly	1.0	Android	11	10	
5	2013-01-01	Acme	1.0	Macintosh	384	214	
6	2013-01-01	Sortly	1.0	Android	14	10	
7	2013-01-01	Sortly	1.0	Windows	1	0	
8	2013-01-01	Acme	0.0	Linux	41	27	
9	2013-01-01	Acme	0.0	iPhone	448	368	
10	2013-01-01	Widgetry	1.0	iPhone	15	14	
11	2013-01-01	Acme	NaN	Windows	58192	46312	
12	2013-01-01	Sortly	1.0	Other	0	0	
13	2013-01-01	Tabular	1.0	iPad	22	21	
14	2013-01-01	Sortly	1.0	Macintosh	0	0	
15	2013-01-01	Acme	NaN	Android	4942	4290	
16	2013-01-01	Acme	NaN	BlackBerry	111	94	
17	2013-01-01	Pinnacle	NaN	Linux	7	4	
18	2013-01-01	Pinnacle	0.0	Macintosh	56	45	
19	2013-01-01	Botly	0.0	Android	20	19	
20	2013-01-01	Widgetry	0.0	iPhone	113	109	
21	2013-01-01	Pinnacle	NaN	BlackBerry	4	2	
22	2013-01-01	Acme	NaN	Linux	984	886	
23	2013-01-01	Pinnacle	NaN	iPhone	282	239	
24	2013-01-01	Acme	NaN	SymbianOS	4	3	

	orders	gross_sales	bounces	add_to_cart	product_page_views	\
0	14	1287.0	4	16	104	
1	0	13.0	0	0	1	
2	0	98.0	0	0	0	
3	527	60753.0	149	610	3914	
4	11	1090.0	0	11	4	
5	213	28129.0	65	245	1783	
6	4	432.0	4	7	33	
7	0	31.0	0	0	2	
8	6	705.0	6	12	130	
9	36	4637.0	80	79	722	
10	15	1813.0	0	15	230	
11	0	NaN	23664	2285	104651	
12	0	NaN	0	0	4	
13	22	3378.0	0	22	0	
14	0	13.0	0	0	3	
15	0	NaN	1751	110	5750	
16	0	NaN	55	3	111	
17	0	NaN	5	0	8	

18	5	1563.0	17	11	120
19	20	2405.0	0	20	0
20	113	15320.0	0	113	1729
21	0	NaN	2	0	4
22	0	NaN	777	10	1078
23	0	NaN	155	6	292
24	0	NaN	2	0	3

	search_page_views	conversion_rate	bounce_rate	add_to_cart_rate
0	192	0.583333	0.166667	0.666667
1	0			
2	0			
3	7367	0.571584	0.161605	0.661605
4	19	1	0	1
5	3255	0.554688	0.169271	0.638021
6	52	0.285714	0.285714	0.5
7	2	0	0	0
8	272	0.146341	0.146341	0.292683
9	1073	0.0803571	0.178571	0.176339
10	216	1	0	1
11	258511	0	0.406654	0.0392666
12	2			
13	0	1	0	1
14	1			
15	14185	0	0.35431	0.0222582
16	198	0	0.495495	0.027027
17	10	0	0.714286	0
18	232	0.0892857	0.303571	0.196429
19	0	1	0	1
20	1466	1	0	1
21	4	0	0.5	0
22	1840	0	0.789634	0.0101626
23	297	0	0.549645	0.0212766
24	4	0	0.5	0

In [43]: mydata.describe() # Summary of numerical variables

```
Out[43]:
```

	new_customer	visits	distinct_sessions	orders \
count	12802.000000	21061.000000	21061.000000	21061.000000
mean	0.448055	1934.708039	1515.205024	62.378994
std	0.497314	7448.607191	5925.833287	260.279286
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	3.000000	2.000000	0.000000
50%	0.000000	24.000000	19.000000	0.000000
75%	1.000000	360.000000	274.000000	7.000000
max	1.000000	136057.000000	107104.000000	4916.000000

	gross_sales	bounces	add_to_cart	product_page_views \
--	-------------	---------	-------------	----------------------

count	11485.000000	21061.000000	21061.000000	21061.000000
mean	16473.395821	743.282085	166.250890	4358.198234
std	51111.354605	3154.697787	505.186834	14327.287354
min	1.000000	0.000000	0.000000	0.000000
25%	79.000000	0.000000	0.000000	3.000000
50%	851.000000	5.000000	4.000000	53.000000
75%	3145.000000	97.000000	43.000000	708.000000
max	707642.000000	54512.000000	7924.000000	187601.000000


	search_page_views
count	21061.000000
mean	8584.187788
std	31120.321365
min	0.000000
25%	4.000000
50%	82.000000
75%	1229.000000
max	506629.000000

```
In [44]: mydata['new_customer'].value_counts()
```

```
Out[44]: 0.0    7066
         1.0    5736
         Name: new_customer, dtype: int64
```

```
In [45]: mydata['product_page_views'].hist()
         plt.xlabel("Number of Page Views per Transaction")
         plt.ylabel("Frequency")
```

```
Out[45]: Text(0,0.5,'Frequency')
```




As we can see in the above histogram there are a few extreme values or outliers. Plotting such distributions will help us in understanding how these outliers can affect the statistical assumptions and analysis of our data. Next, let us take a look at boxplots of some of columns in **mydata**.

Below are two boxplots of the *search_page_views* column with extreme values and without extreme values. The first plot shows a heavy presence of outliers, which makes our distribution very skewed. From the second plot, we can see that the minimum value is the same as the 1st (lower) quartile.


```
In [46]: mydata.boxplot(column='search_page_views')
         plt.ylabel("Frequency")
```

```
Out[46]: Text(0,0.5, 'Frequency')
```



```
In [47]: mydata.boxplot(column='search_page_views', showfliers=False)
plt.ylabel("Frequency")
```

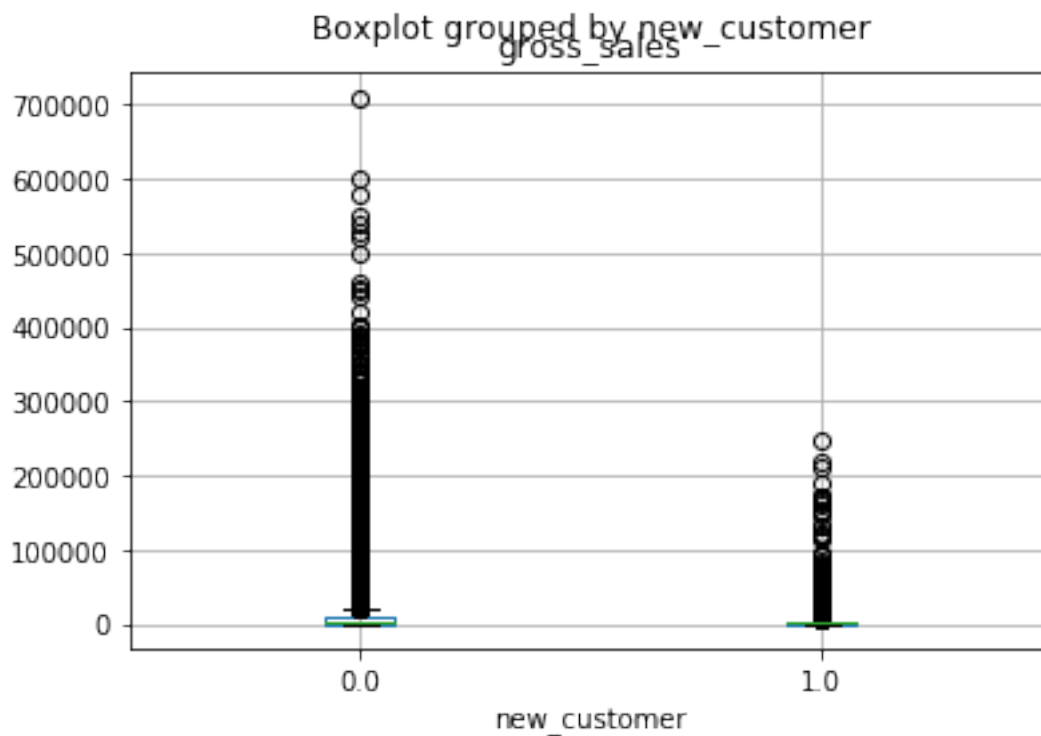
```
Out[47]: Text(0,0.5, 'Frequency')
```



One question we may ask is what is the comparison between the gross sales made by new customers (1) versus returning customers (0). Using the below two boxplots, we are able to see that the median and range of gross sales for returning customers is quite bigger than those of new customers. However, both plots indicate that the distribution of the gross sales for each type of customer dimension is skewed right. The first boxplot shows countless outliers that make it nearly impossible to see the spread of the data column. By removing the outliers, we get a better look of the distribution (second plot).

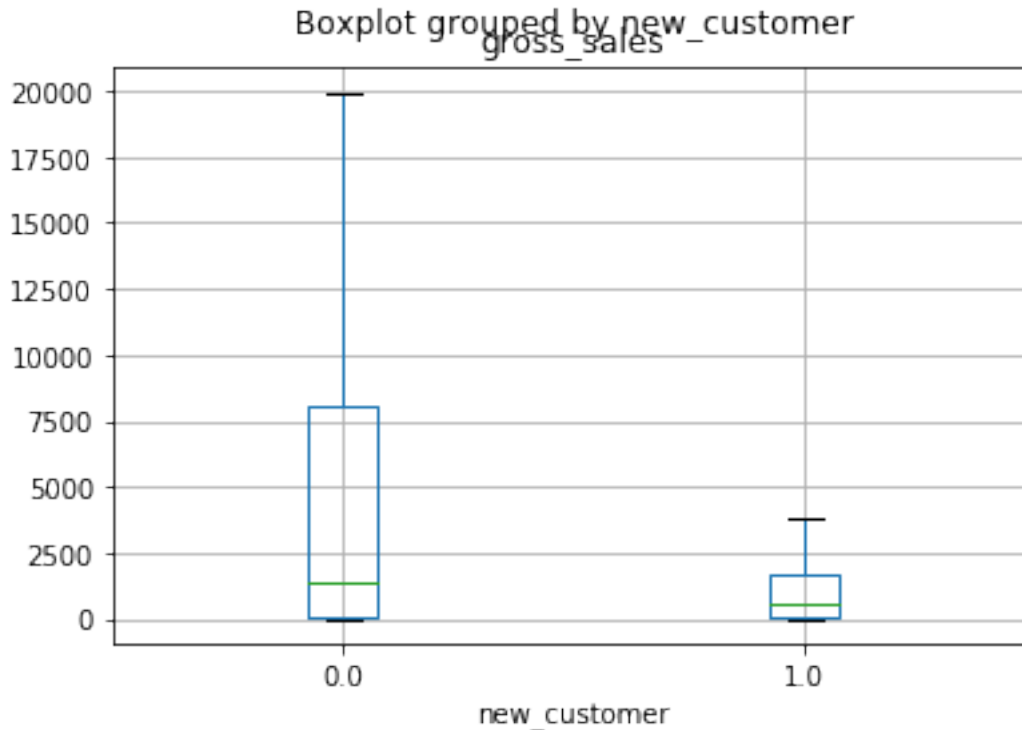
```
In [48]: mydata.boxplot(column='gross_sales', by= 'new_customer')
```

```
Out[48]: <matplotlib.axes._subplots.AxesSubplot at 0x13f7cef6668>
```



```
In [49]: mydata.boxplot(column='gross_sales', by= 'new_customer', showfliers=False)
```

```
Out[49]: <matplotlib.axes._subplots.AxesSubplot at 0x13f7be0f588>
```



We can see how strong the relationship between the types of customers we've had and gross sales by finding the correlation coefficient, which will solidfy the evidence we saw in the distribution plots. The results shows that there isn't a strong nor positive relationship between the users and the gross sales.

```
In [50]: mydata['new_customer'].corr(mydata['gross_sales'])
```

```
Out[50]: -0.21188490786384043
```

We want to see the sales, orders, visits and bounces by month. First, we will have to convert *day* to be datetime format using the **datetime** package.

```
In [51]: import datetime
```

```
mydata['year'] = pd.DatetimeIndex(mydata['day']).year
mydata['month'] = pd.DatetimeIndex(mydata['day']).month
mydata['dow'] = pd.DatetimeIndex(mydata['day']).weekday_name
mydata['myday'] = pd.DatetimeIndex(mydata['day']).day

print(mydata['year'])
print(mydata['month'])
print(mydata['dow'])
print(mydata['myday'])
```


0	2013
1	2013
2	2013
3	2013
4	2013
5	2013
6	2013
7	2013
8	2013
9	2013
10	2013
11	2013
12	2013
13	2013
14	2013
15	2013
16	2013
17	2013
18	2013
19	2013
20	2013
21	2013
22	2013
23	2013
24	2013
25	2013
26	2013
27	2013
28	2013
29	2013
	...
21031	2013
21032	2013
21033	2013
21034	2013
21035	2013
21036	2013
21037	2013
21038	2013
21039	2013
21040	2013
21041	2013
21042	2013
21043	2013
21044	2013
21045	2013
21046	2013
21047	2013

21048	2013
21049	2013
21050	2013
21051	2013
21052	2013
21053	2013
21054	2013
21055	2013
21056	2013
21057	2013
21058	2013
21059	2013
21060	2013

Name: year, Length: 21061, dtype: int64

0	1
1	1
2	1
3	1
4	1
5	1
6	1
7	1
8	1
9	1
10	1
11	1
12	1
13	1
14	1
15	1
16	1
17	1
18	1
19	1
20	1
21	1
22	1
23	1
24	1
25	1
26	1
27	1
28	1
29	1
	..
21031	12
21032	12
21033	12

21034	12
21035	12
21036	12
21037	12
21038	12
21039	12
21040	12
21041	12
21042	12
21043	12
21044	12
21045	12
21046	12
21047	12
21048	12
21049	12
21050	12
21051	12
21052	12
21053	12
21054	12
21055	12
21056	12
21057	12
21058	12
21059	12
21060	12

Name: month, Length: 21061, dtype: int64

0	Tuesday
1	Tuesday
2	Tuesday
3	Tuesday
4	Tuesday
5	Tuesday
6	Tuesday
7	Tuesday
8	Tuesday
9	Tuesday
10	Tuesday
11	Tuesday
12	Tuesday
13	Tuesday
14	Tuesday
15	Tuesday
16	Tuesday
17	Tuesday
18	Tuesday
19	Tuesday

```

20      Tuesday
21      Tuesday
22      Tuesday
23      Tuesday
24      Tuesday
25      Tuesday
26      Tuesday
27      Tuesday
28      Tuesday
29      Tuesday
      ...
21031   Tuesday
21032   Tuesday
21033   Tuesday
21034   Tuesday
21035   Tuesday
21036   Tuesday
21037   Tuesday
21038   Tuesday
21039   Tuesday
21040   Tuesday
21041   Tuesday
21042   Tuesday
21043   Tuesday
21044   Tuesday
21045   Tuesday
21046   Tuesday
21047   Tuesday
21048   Tuesday
21049   Tuesday
21050   Tuesday
21051   Tuesday
21052   Tuesday
21053   Tuesday
21054   Tuesday
21055   Tuesday
21056   Tuesday
21057   Tuesday
21058   Tuesday
21059   Tuesday
21060   Tuesday
Name: dow, Length: 21061, dtype: object
0      1
1      1
2      1
3      1
4      1
5      1

```

6	1
7	1
8	1
9	1
10	1
11	1
12	1
13	1
14	1
15	1
16	1
17	1
18	1
19	1
20	1
21	1
22	1
23	1
24	1
25	1
26	1
27	1
28	1
29	1
	..
21031	31
21032	31
21033	31
21034	31
21035	31
21036	31
21037	31
21038	31
21039	31
21040	31
21041	31
21042	31
21043	31
21044	31
21045	31
21046	31
21047	31
21048	31
21049	31
21050	31
21051	31
21052	31
21053	31

```

21054    31
21055    31
21056    31
21057    31
21058    31
21059    31
21060    31
Name: myday, Length: 21061, dtype: int64

```

After creating new columns **year**, **month**, **dow** to hold the year, month and day of the week (Monday-Sunday), we can find the distributions grouped by these columns.

Sales by Month

```

In [52]: #Sales
         sales_by_month = mydata.groupby(mydata['month']).size()
         print(sales_by_month)

         #Plotting the graph
         plot_by_month = sales_by_month.plot(title='Monthly Sales', xticks=(1,2,6,7,8,9,10,11,12))
         plot_by_month.set_xlabel('Months')
         plot_by_month.set_ylabel('Number of Sales')

```

```

month
1      2366
2      2137
6      2327
7      2035
8      2462
9      2347
10     2464
11     2389
12     2534
dtype: int64

```

```

Out[52]: Text(0,0.5,'Number of Sales')

```



Sales by Day

```
In [53]: #Sales
sales_by_day = mydata.groupby(mydata['myday']).size()
print(sales_by_day)

#Plotting the graph
sales_by_day = sales_by_day.plot(title='Daily Sales', xticks=(range(0,31)), rot=55)
sales_by_day.set_xlabel('Day')
sales_by_day.set_ylabel('Number of Sales')
```


```
myday
1      704
2      710
3      713
4      697
5      709
6      698
7      712
8      704
9      697
10     630
11     627
12     626
13     624
```

```

14    627
15    716
16    710
17    715
18    706
19    712
20    701
21    697
22    709
23    712
24    709
25    711
26    712
27    705
28    708
29    633
30    637
31    390
dtype: int64

```

```
Out[53]: Text(0,0.5,'Number of Sales')
```



Sales by Day of the Week


```
In [54]: #Sales
sales_by_dow = mydata.groupby(mydata['dow']).size()
print(sales_by_dow)

#Plotting the graph
sales_by_dow = sales_by_dow.plot(title='Day of Week Sales', xticks = range(0,7), rot=
sales_by_dow.set_xlabel('Day of Week')
sales_by_dow.set_ylabel('Number of Sales')
```

```
dow
Friday      2895
Monday      3064
Saturday     2973
Sunday       2971
Thursday     3006
Tuesday      3151
Wednesday   3001
dtype: int64
```

```
Out[54]: Text(0,0.5,'Number of Sales')
```

